Voyage AI

Embedding models

(RAG) is the Predominant Approach for Enterprise Gen AI

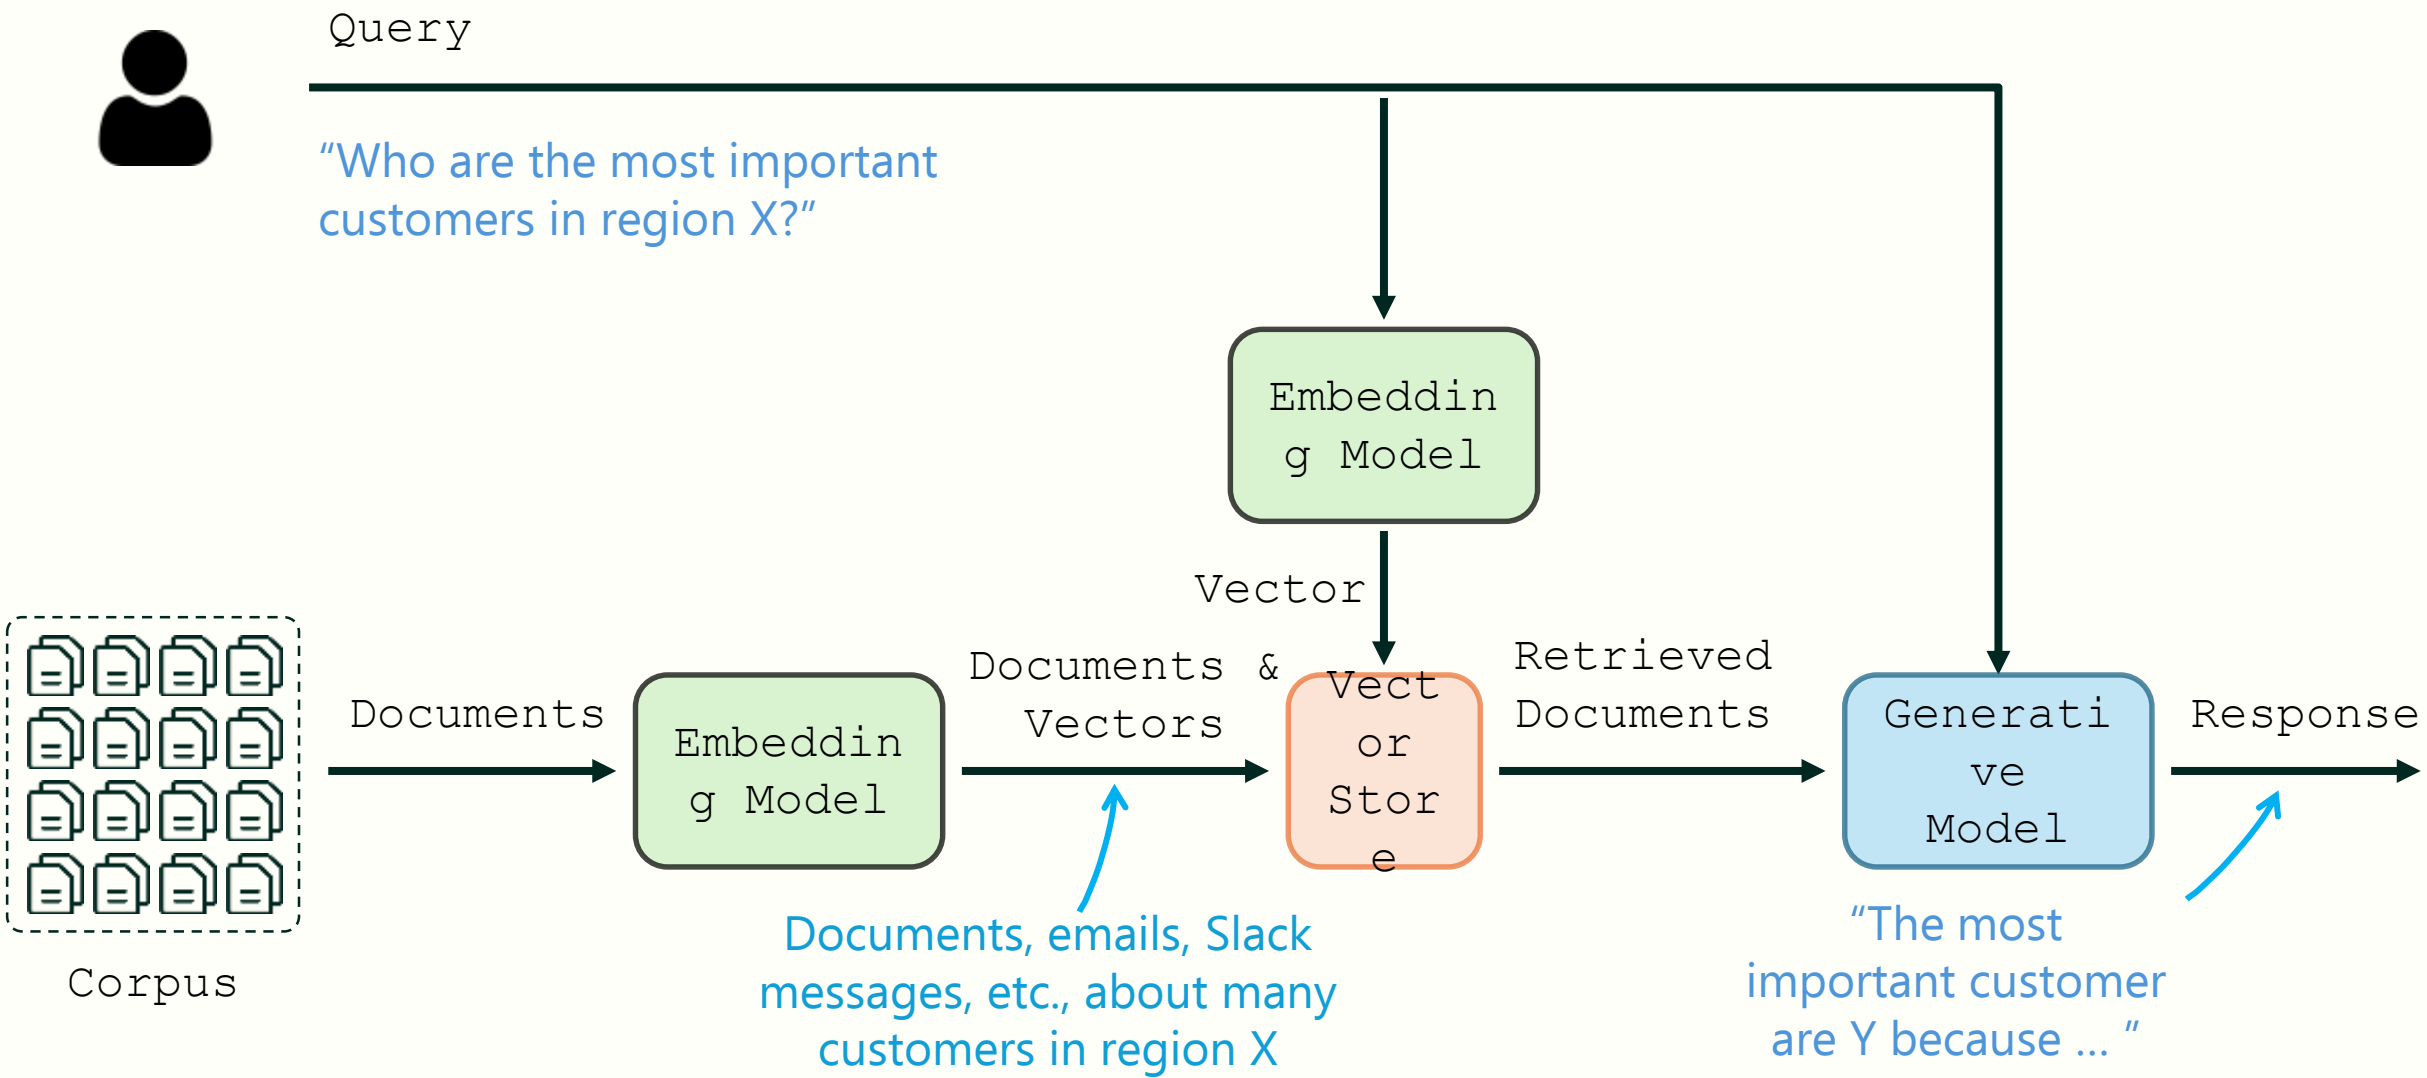

(RAG) is the Predominant Approach for Enterprise Gen AI

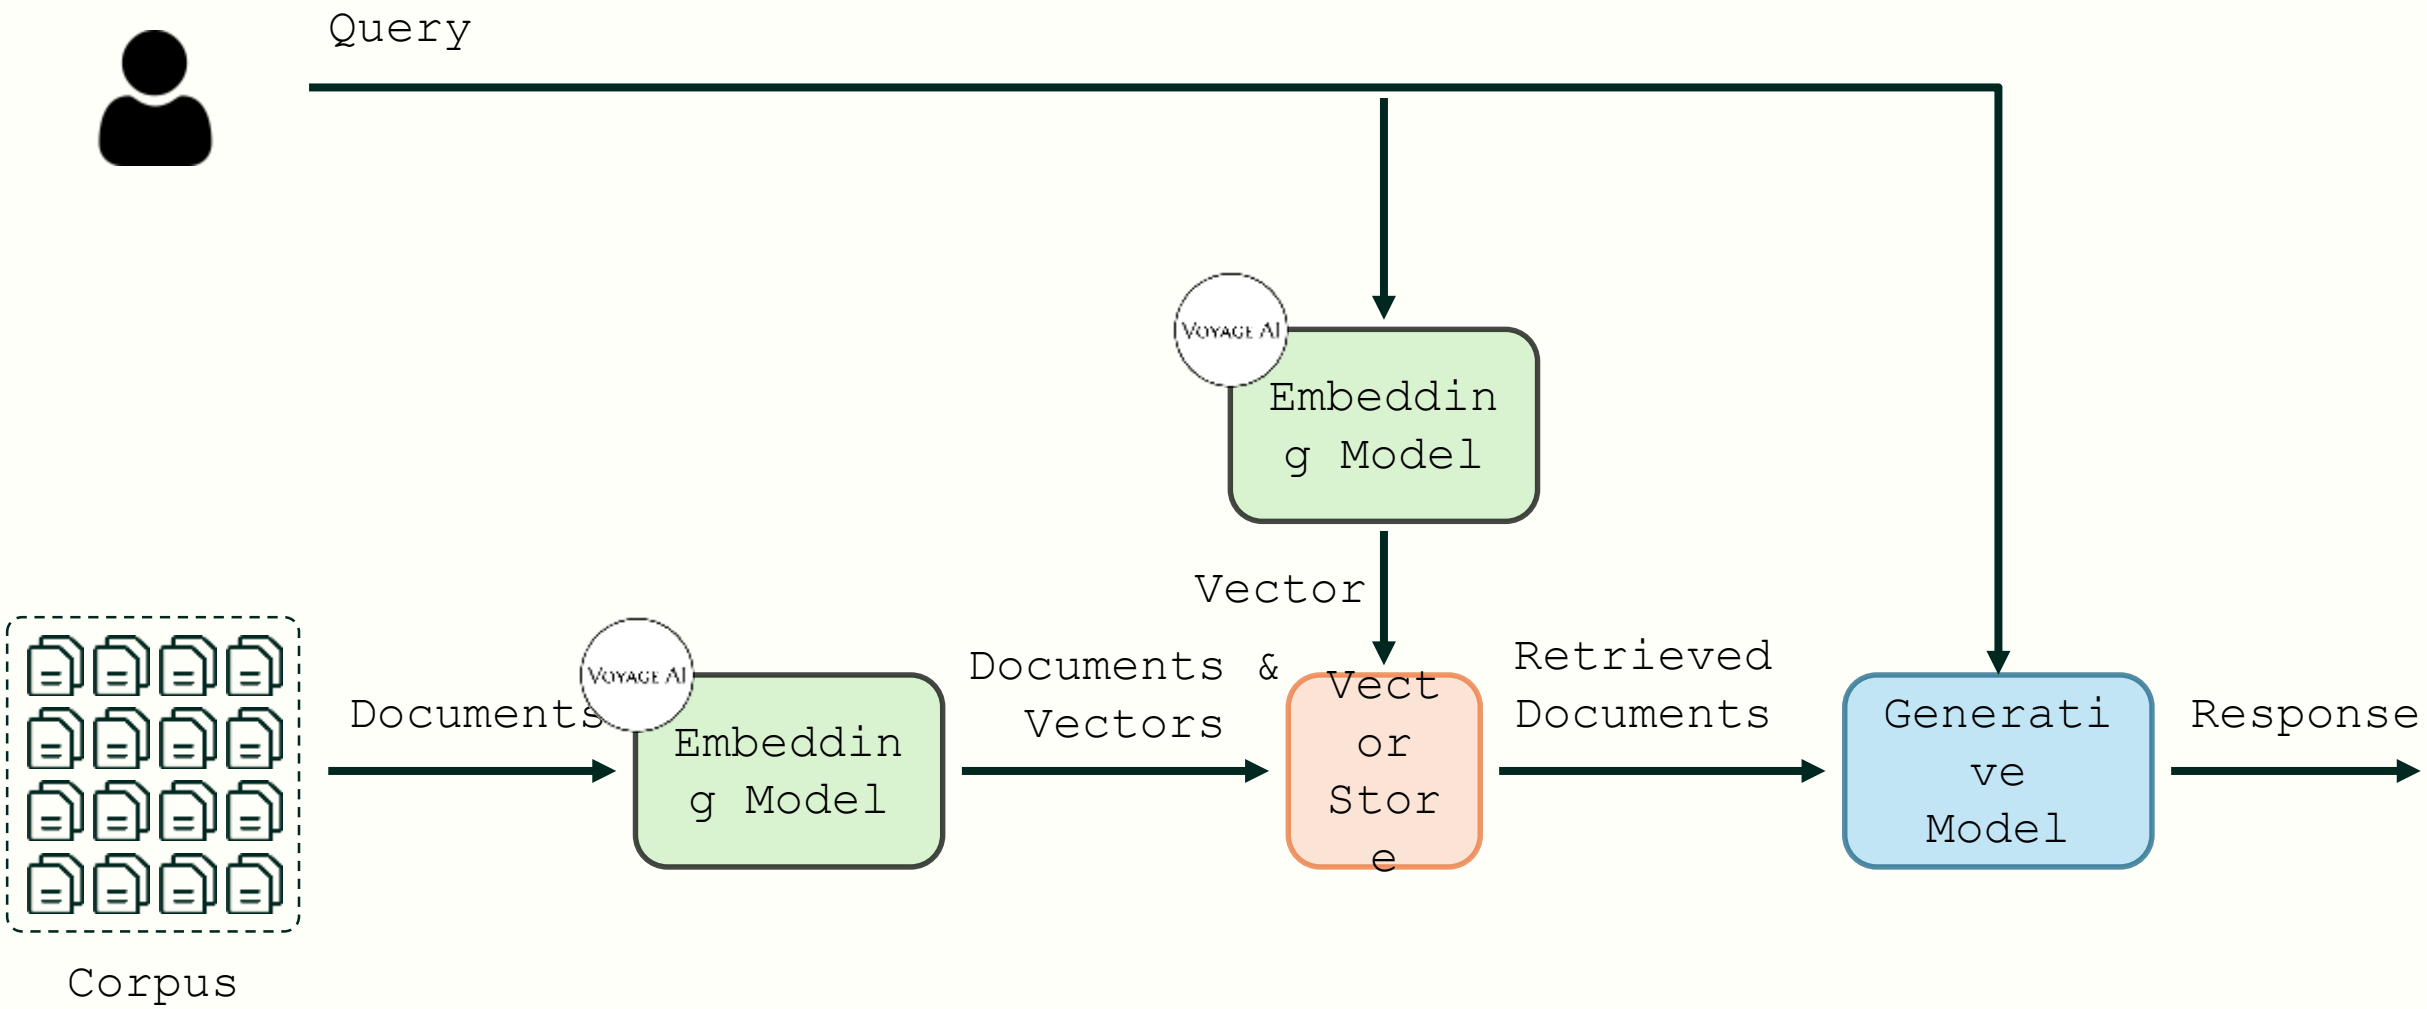

Voyage AI Products and Offerings

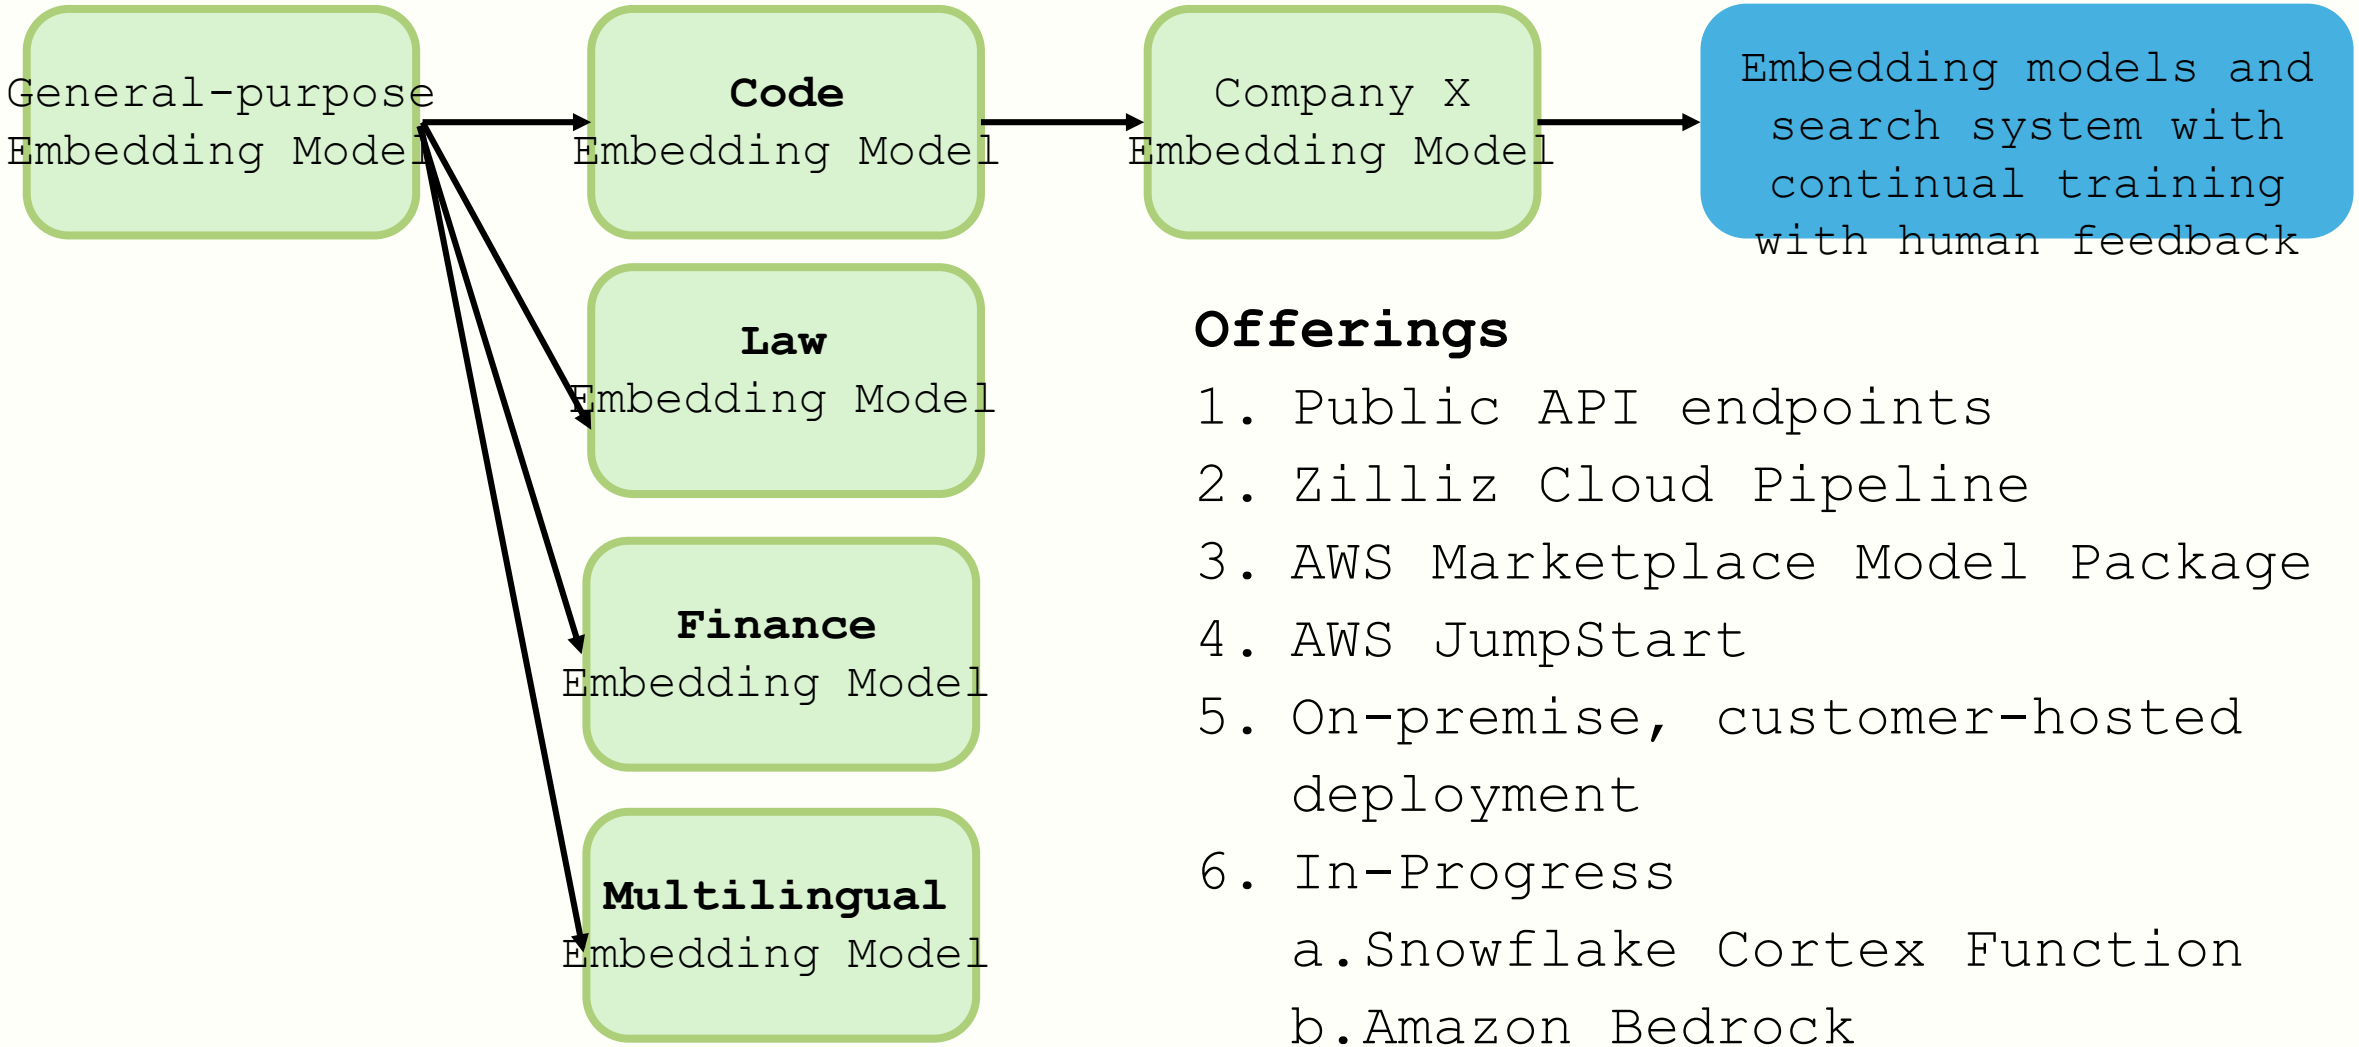

Offerings

1. Public API endpoints
2. Zilliz Cloud Pipeline
3. AWS Marketplace Model Package
4. AWS JumpStart
5. On-premise, customer-hosted deployment
6. In-Progress
   a. Snowflake Cortex Function
   b. Amazon Bedrock

Voyage Models are State-of-the-Art in Retrieval Quality

Recall@5 for Industry-domain Retrieval Datasets

Models \ Datasets	Industry-Average	Langchain Docs	Health4CA Policies	DoorDash Reviews	PyTorch Docs	CNN Sports News	Cohere Docs	Verizon 5G Docs	Huffpost Science News	Onesignal Docs
BAAI/bge	77.67	63.19	88.83	54.19	76.41	89.00	85.75	84.89	92.61	64.20
Cohere v3	82.03	65.88	90.08	58.77	84.34	95.85	86.28	87.11	96.18	73.78
OpenAI-3-large	83.93	64.70	92.11	51.29	90.25	**97.73**	88.14	92.79	97.55	80.85
voyage-2	86.68	**74.20**	**92.18**	59.42	93.78	96.00	91.25	**93.74**	96.62	82.91
voyage-large-2	**86.96**	73.19	92.05	**59.55**	**94.42**	97.11	**91.74**	93.01	**97.90**	**83.68**

The Code Embedding Model Excels on Code Retrieval Tasks

Recall@5 for Code Retrieval Datasets

Models \ Datasets	Code-Average	APPS	HumanEval-Text	HumanEval-TextCode	CodeChefCPP	CodeChef Python	MBPP	DS-1000	DS-1000-ReferenceOnly	LeetCodeCpp	LeetCodeJava	LeetCode Python
BAAI/bge	53.74	5.93	73.42	78.66	4.74	9.24	88.40	46.40	44.65	75.87	80.11	83.73
Cohere v3	62.68	11.56	87.98	90.24	8.82	18.18	88.50	58.61	65.57	85.04	88.10	86.90
OpenAI-3-large	75.84	24.45	**99.37**	**100.00**	25.61	29.13	94.56	80.48	84.18	98.89	98.91	98.70
voyage-code-2	**90.36**	**74.17**	**99.37**	**100.00**	**72.17**	**76.00**	**96.41**	**88.84**	**87.24**	**100.00**	**99.91**	**99.90**

Early Pilot Partners

Trusted by Partners like ANTHROP\C

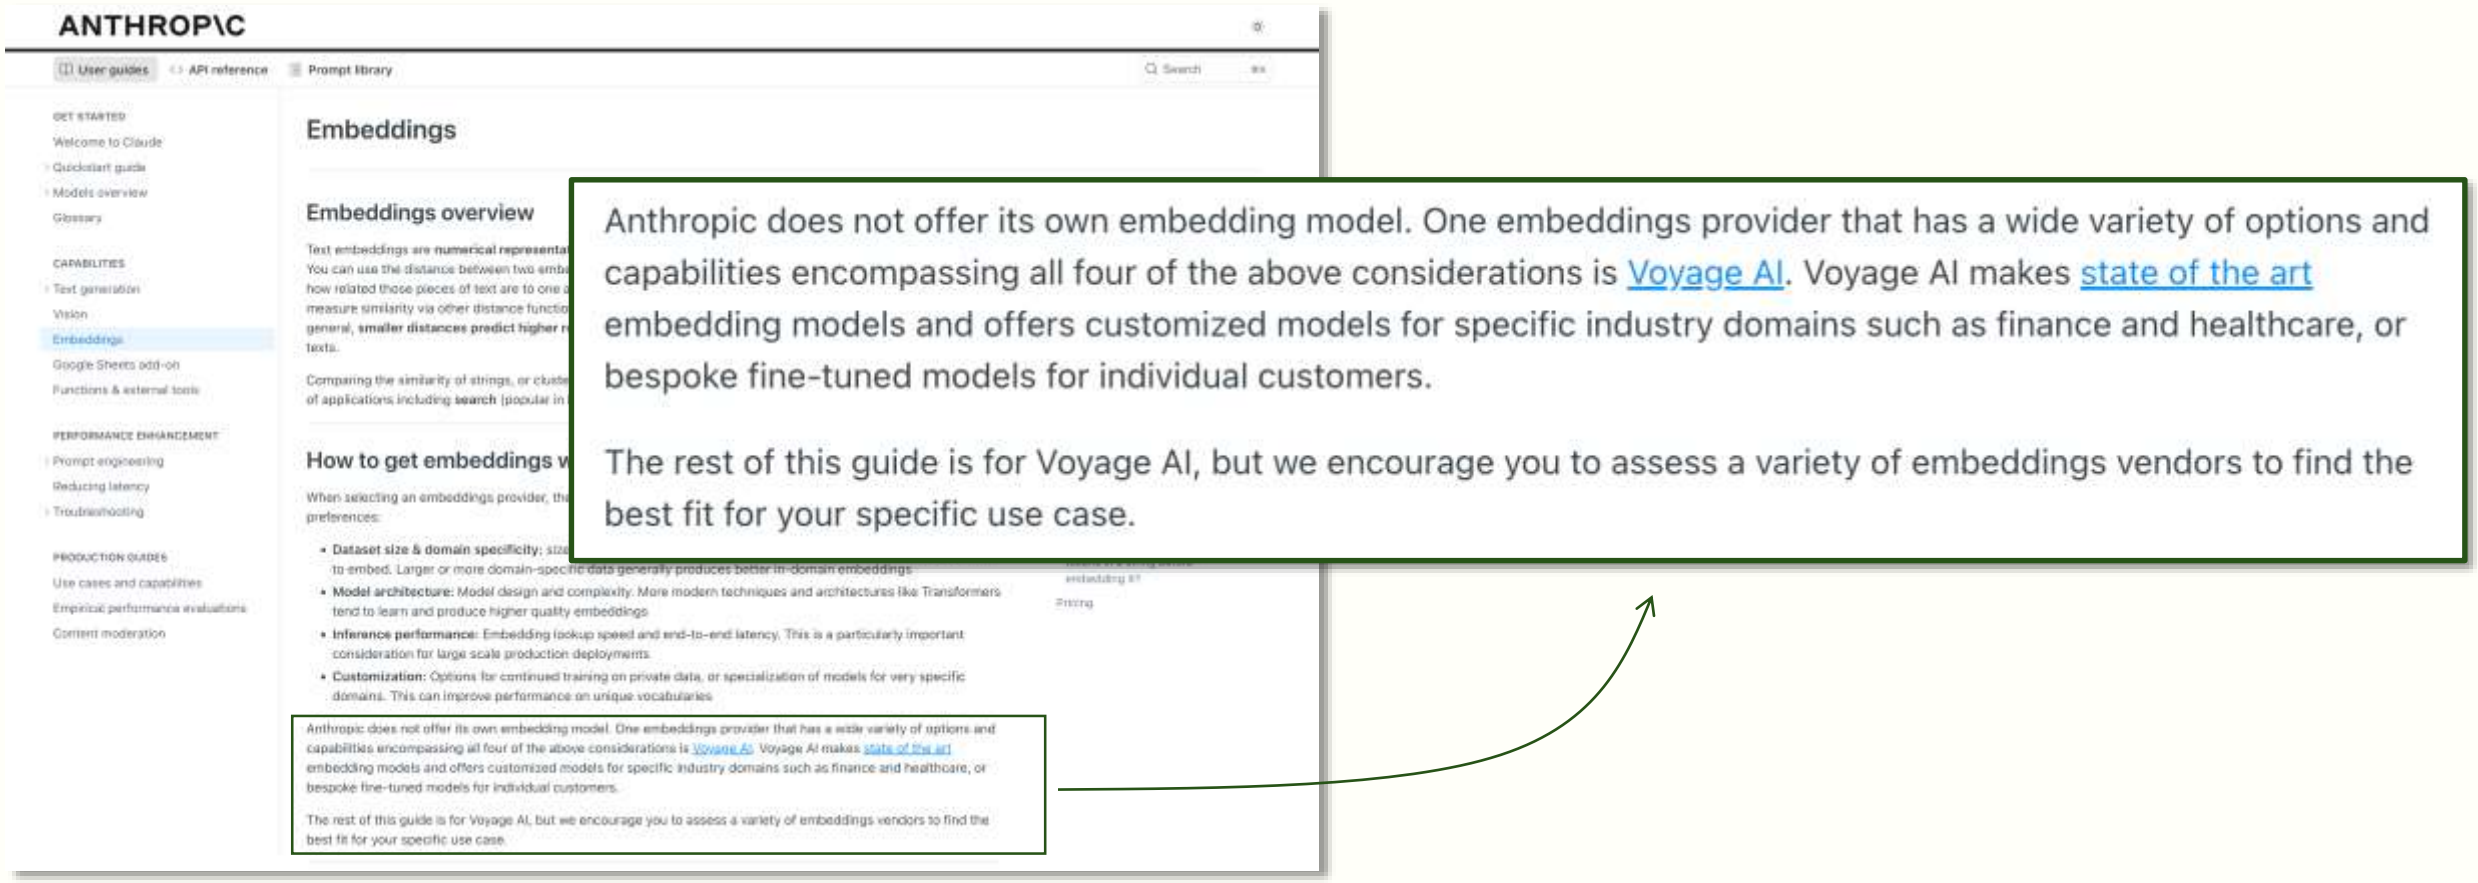

Trusted by Zilliz

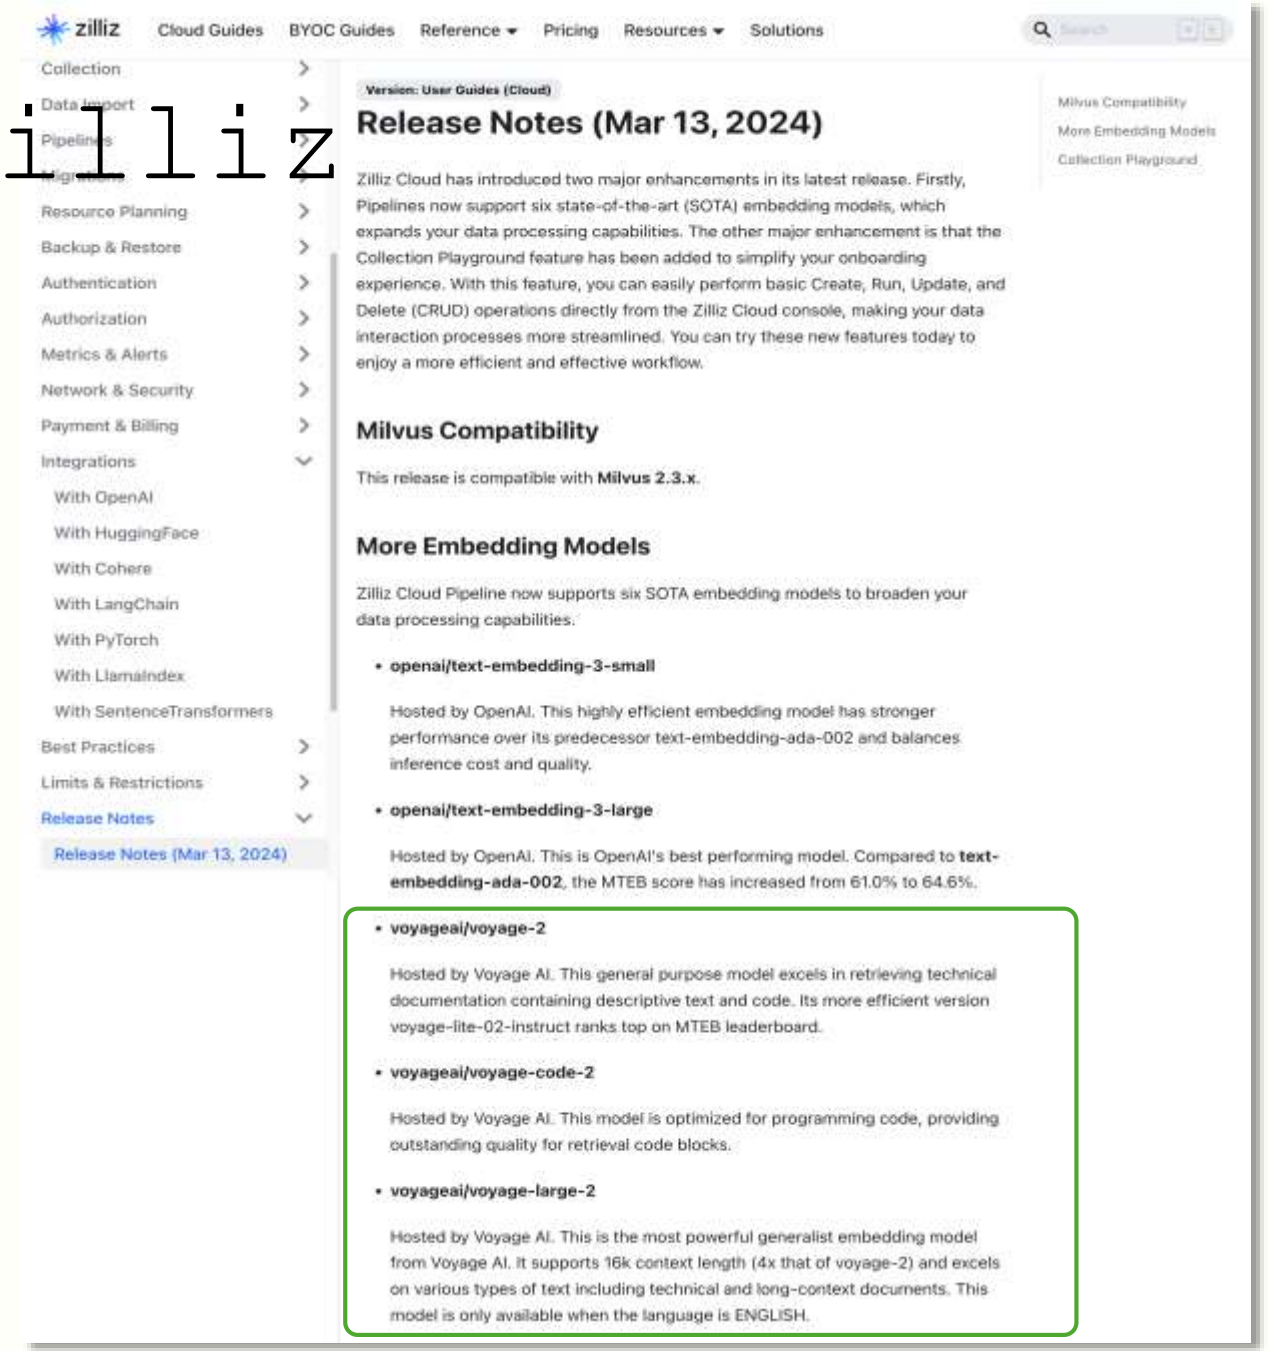

Use Voyage Embeddings on Zilliz Cloud Pipeline